

LMS Code Walk

TIARK’S GROUP

Lightweight Modular Staging (LMS) is a generative programming tool achieving the multi-stage programming (staging). LMS_clean is a remake of the LMS project, aiming at a more flexible design and extension with better support for LMS IR transformation and optimization. This documentation is a code walk of the LMS_clean repo, hoping to explain the implementation (in high-level or in details) to people who are interested in learning and using LMS. This is different from a tutorial since it will dive into the core implementation of LMS and offers more insights than simply how to use LMS.

1 INTRODUCTION

Multi-Stage Programming (or staging) is a programming language concept that allows various parts of the programs to run in different stages. It allows users to code with high-level abstractions (trait, classes, high-order functions), but still gain highly-efficient code after the abstractions are executed (staged away). This offers both high productivity and high performance of the target program, thus the slogan “abstract without regret”.

Lightweight Modular Staging (LMS) is a staging tool built in Scala. In LMS, type information is used to distinguish the evaluation stages (i.e., All $\text{Rep}[T]$ typed values and expressions are in code generation for the next stage.) Simply speaking, LMS is a compiler. However, LMS does not have a laxer or parser to transform the input program into intermediate representations (IR). Instead, the IR is generated via executing the input program. All the $\text{Rep}[T]$ typed expressions in the input program evaluate to LMS IR. This can be considered as the LMS frontend. Then the LMS backend compiles the LMS IR to target programs.

2 LMS IR

In the file core/backend.scala, object Backend, the core LMS IR is described.

```
abstract class Def // Definition: used in right-hand-side of all nodes
abstract class Exp extends Def
case class Sym(n: Int) extends Exp // Symbol
case class Const(x: Any) extends Exp // Constant
case class Block(in: List[Sym], res: Exp, ein: Sym, eff: EffectSummary) extends Def
case class Node(n: Sym, op: String, rhs: List[Def], eff: EffectSummary)
```

2.1 Sea Of Nodes

The LMS IR follows the “sea of nodes” design (cite?), where the IR is composed of a list of Nodes, and the Blocks do not explicitly scope the nodes. Instead, the Blocks describe their *inputs* (via *in*: List[Sym]), *result* (via *res*: Exp), *world* (via *ein*: Sym), and *effects* (via *eff*: EffectSummary).

The *world* (by *ein*: Sym) is probably unfamiliar to many readers. It is used to scope the nodes that belong to this block. That is to say, if a node depends on the world of a block, then it should be scheduled within that block. For example:

Author’s address: Tiark’s Group.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

© 2020 Association for Computing Machinery.

XXXX-XXXX/2020/6-ART \$15.00

<https://doi.org/10.1145/nnnnnnnn.nnnnnnnn>

Give an example?

If a node depends on the inputs of a block, very likely it has to be scheduled in the block too. So both inputs and world are scoping the beginning of the block. That is why the world is also called the “effect input”. So if the world marks the beginning of a block, which component of the Block marks the end of a block? The answer is the result and the effects of the block. When scheduling the nodes for a block, we can start from the result and effects, and pick nodes that are depended on by them.

Why using the sea of node IR in LMS? One of the key reasons is that it offers easy optimization of code-motion. Since the blocks do not explicitly scope the nodes, whether a node is scheduled in or out of a block is flexibly determined based on effects, using frequency, and et al. It does make IR traversal and transformation different from IRs with explicit scopes, which we will talk about in `core/traversal.scala`.

2.2 Effect Summary

In Section 2.1, we talked about the dependencies between nodes and blocks (via `EffectSummary`). `EffectSummary` can be course-grained or fine-grained. More fine-grained effect summary will have higher compilation time cost, but provide more optimized target program.

First of all, we want to be sure of the two intertwined concepts here: effects and dependencies. Effects refer to node behaviors such as printing, variable read, variable write, et al. Dependencies refer to both data dependencies and dependencies caused by effects, such as printing cannot be missed or reordered, and we cannot reorder two writes on the same variable.

The general work flow is that at node constructions, the effects of nodes are tracked. using the node effects we compute the dependencies. The dependencies are then used to schedule nodes for each block. Our current `EffectSummary` is like this:

```
case class EffectSummary(sdeps: Set[Sym], hdeps: Set[Sym], rkeys: Set[Exp], wkeys: Set[Exp])
```

It tracks soft dependencies (via `sdeps: Set[Sym]`), hard dependencies (via `hdeps: Set[Sym]`), read keys (via `Set[Exp]`), and write keys (via `Set[Exp]`). Hard dependencies are the normal dependencies. If node A hard-depends on node B, then scheduling A ensures that B is scheduled before A. Soft dependencies are soft in the sense that, if node A soft-depends on node B, node B cannot be scheduled after A. However, scheduling A does not ensure that B is scheduled. Read keys and write keys are easy to understand: they track Exps that are read or written to by the node or block. However, the read keys and write keys can also be `Const`, indicating that they have other semantics related to other types of effects. For instance, printing nodes have write effect on `Const("CTRL")`, which ensures that printings are all scheduled with the order unchanged. Allocating variables/arrays have read effects on `Const("Store")`.

3 LMS GRAPH

In this section, we will talk about how LMS Graph are constructed using the LMS IR components. It shows how the LMS IRs are used in constructing LMS Graphs, and how effects and dependencies are tracked and generated. All LMS snippets are function. As the result, all LMS Graph have a list of nodes (already in topological order) and a block describing the function. That is captured by the `case class` `Graph(val nodes: Seq[Node], val block: Block, val globalDefsCache: immutable.Map[Sym, No` at `core/backend.scala`. The LMS Graph is constructed by class `GraphBuilder` at `core/backend.scala`.

Besides the basic functionality of storing nodes, searching nodes by symbols, and generating fresh symbols, `GraphBuilder` offers two keys functionalities

- (1) Building nodes by the `reflect*` family of methods.

(2) Building blocks by the `reify*` family of methods.

Note that the `reify*` family of methods not only generate the `Block` object, but also the nodes that are used in the block. However, the nodes are not explicitly scoped in the block, but rather implicitly scoped via effect summaries. This implicit scoping allows flexible code motion as long as effects and dependencies are respected.

3.1 High-level Design of the Effect System

3.2 Code Walk: `reflectEffect`

The core method in the `reflect*` family of methods is the `reflectEffect` method.

```
def reflectEffect(s: String, as: Def*)(readEfKeys: Exp*)(writeEfKeys: Exp*): Exp
```

3.3 Code Walk: `reify`

4 SIMPLE FRONTEND

With a proper definition of LMS IR and the facility to build LMS Graph in `core/backend.scala`, we can build a frontend that construct LMS Graph. The LMS frontend should feature the `Rep[T]` type, which allows the staged programs to be type checked. However, in this section, we are going to introduce a very simple frontend that can just nit the LMS Graph, without the iconic `Rep[T]` type. This simple frontend is not of much use in production, but it shows the simple essence of LMS front, i.e., being able to construct LMS Graph with various control flows.

The basic ways to construct LMS graphs are through the `reflect*` and `reify*` family of functions. However, the frontend should allow the users to construct LMS graphs via unary and binary operations, `If`, `While`, `FUN`, and et al. That is the main purpose of the simple `FrontEnd` class.

```
class FrontEnd {
  var g: GraphBuilder = null // LMS graph is built in here

  case class BOOL(x: Exp) // boolean wrapper of LMS EXP, supporting ! op
  case class INT(x: Exp) // int-like wrapper of LMS EXP, supporting arithmetic op
  case class ARRAY(x: Exp) // array wrapper of LMS Exp, supporting array access.
  case class VAR(x: Exp) // variable wrapper of LMS Exp, supporting variables

  // supporting conditional. a and b are executed in reify*.
  // then the returned blocks are used in reflect* to build the conditional node
  def IF(c: BOOL)(a: => INT)(b: => INT): INT = {...}

  // supporting loop. both c and b are executed in reify*,
  // then the returned blocks are used in reflect* to build the loop node
  def WHILE(c: => BOOL)(b: => Unit): Unit = {...}

  // supporting application. just creating an app node.
  def APP(f: Exp, x: INT): INT = INT(g.reflect("@", f, x.x))

  // If we just need to support in-graph, non-recursive functions, this is enough.
  // It builds a lambda node with the `f` reified into the block,
  // then it returns a scala function that can be applied to create APP construct
  def FUN(f: INT => INT): INT => INT = {
    val fn = INT(g.reflect("lambda", g.reify {xn: Exp => f(INT(xn)).x} ) )
    (x: INT) => APP(fn.x, x)
  }
}
```

It would be interesting to find out how to achieve in-graph, recursive functions. If we want to support recursive functions, we have to create a lambda node that uses the lambda `Exp` within the lambda block.

```

148 val fn = Sym(g.fresh)
149 g.reflect(fn, "lambda", g.reify(???))

```

In order to be able to use the same `fn` in the block of the `lambda`, we need to be able to construct `APP(fn, xn)` within the `g.reify`.

```

152 val f1 = (x: INT) => APP(fn, x)

```

But we don't know how the `f1` is recursively used. That has to come from user code `f`.

```

153 def FUN(f) = {
154   ...
155   g.reflect(fn,"lambda",g.reify(xn: Exp => f(f1, INT(xn)).x))()
156 }
157 // user code `f` decides how `f1` (the in-graph lambda) is recursively used.

```

To put everything together and offer a non-recursive API:

```

158 def FUN(f: INT => INT): INT => INT = FUN( (_,x) => f(x))
159
160 def FUN(f: ((INT=>INT),INT) => INT): INT => INT = {
161   val fn = Sym(g.fresh)
162   val f1 = (x: INT) => APP(fn,x)
163   g.reflect(fn,"lambda",g.reify(xn => f(f1,INT(xn)).x))()
164   f1
165 }

```

A use case might be

```

166 val fac = FUN { (f, n) =>
167   IF (n != 0) {
168     n * f(n-1) // recursive call
169   } {
170     1
171   }
172 }

```

Finally, the program function reifies the user provided snippet (of type `INT => INT`) and returns the LMS Graph.

5 NORMAL FRONTEND

After discussing the simple frontend in Section 4, we want to show what a normal frontend looks like (in the object `Adaptor` at `core/stub.scala`).

object `Adaptor` extends `Frontend` typeMap funTable

emitCommon // code gen ???

class `MyGraphBuilder` extends `GraphBuilder`

Base `EmbeddedControls` (macro and virtualization) `OverloadHack` (for hacky overload) `ClosureCompare` (compare LMS closures)

5.1 Base: Introducing Rep[T]

In the `Base` trait, the code establish the iconic `Rep[T]` of LMS. Previously in `Frontend` class, we have seen one way to wrap around `core.backend.Exp` so that we can construct LMS Graph via unary operators, binary operators, and et al. What is to be further provided in `Base` trait is the ability to use `Rep[T]`. Similarly `Rep[T]` is built on top of `core.backend.Exp`. The `core.backend.Exp` do not have types. The types are added via user code and type inferencing, and then tracked in a data structure called typeMap (such as `Adaptor.typeMap`).

```

191 trait Base extends EmbeddedControls with OverloadHack with ClosureCompare {
192   type Rep[+T] = Exp[T] // type name aliasing
193
194   abstract class Exp[+T] // track LMS IR for non-variables
195   abstract class Var[T] // track LMS IR for variables
196

```

```

197 // The Wrap class and method that build Rep[T] typed expression with type tracking
198 case class Wrap[+A:Manifest](x: lms.core.Backend.Exp) extends Exp[A] {
199   Adapter.typeMap(x) = manifest[A]
200 }
201 def Wrap[A:Manifest](x: lms.core.Backend.Exp): Exp[A] = {
202   if (manifest[A] == manifest[Unit]) Const().asInstanceOf[Exp[A]]
203   else new Wrap[A](x)
204 }
205 // The WrapV class for Var[T]
206 case class WrapV[A:Manifest](x: lms.core.Backend.Exp) extends Var[A] {
207   Adapter.typeMap(x) = manifest[A]
208 }
209 }
210
211

```

5.2 Base: Better Handling of Functions

In the simple frontend, we see that the handling of recursive functions is a bit awkward. How do we make it better with a better frontend? In Base trait, we express the type of in-graph function better:

```

216 def fun[A:Manifest, B:Manifest](f: Rep[A=>Rep[B]]): Rep[A=>B] =
217   Wrap[A=>B](__fun(f, 1, xn => Unwrap(f(Wrap[A](xn(0)))))

```

Unfortunately, we have to implement multiple funs for different function arities, which we will elide here. The __fun function reifies the argument into a function block.

```

220 def __fun[T:Manifest](f: AnyRef, arity: Int, gf: List[Backend.Exp] => Backend.Exp): Backend.Exp =
221   // use canonicalize to get the unique representation of any Scala function
222   val can = canonicalize(f)
223   Adapter.funTable.find(_._2 == can) match {
224     case Some((funSym, _)) =>
225       funSym // Easy case: found the function in funTable
226     case _ =>
227       // Step 1. set up "lambdaforward" node with 2 new fresh Syms
228       val fn = Backend.Sym(Adapter.g.fresh)
229       val fn1 = Backend.Sym(Adapter.g.fresh)
230       Adapter.g.reflect(fn, "lambdaforward", fn1)()
231
232       // Step 2. register (fn, can) in funTable, so that recursive calls
233       // will find fn as the function Sym. Reify the block.
234       // Note: it might seem strange why/how recursive calls re-enter this __fun() function.
235       // The reason is that in user code, recursive functions have to be written as
236       // lazy val f = fun{...} or def f = fun{...}, in which case the recursive calls
237       // will re-enter the `fun` call.
238       Adapter.funTable = (fn, can)::Adapter.funTable
239       val block = Adapter.g.reify(arity, gf)
240
241       // Step 3. build the "lambda" node with fn1 as the function name
242       // fix the funTable such that it pairs (fn1, can) for non-recursive uses.
243       val res = Adapter.g.reflect(fn1, "lambda", block)(hardSummary(fn))
244       Adapter.funTable = Adapter.funTable.map {
245         case (fn2, can2) => if (can == can2) (fn1, can) else (fn2, can2)
246       }
247       res
248   }
249 }

```